DEVLOOPS CONSULTING · PRODUCTION INFRASTRUCTURE FOR AI & ROBOTICS

We take your AI and robots from demo to dependable production.

Your engineers build the product. We build and run the platform under it: cloud, Kubernetes, GPUs, MLOps and fleet operations.

Ship sooner

From prototype to customers in weeks, not quarters.

Stay up

Systems your customers and investors can trust.

Spend less runway

Lower cloud and GPU cost per model and per robot.

devloopsconsulting.com

THE WORLD TODAY

Money is pouring into AI and robotics. Infrastructure decides who wins.

~50%

of all global venture funding in 2025 went to AI companies.

$18.8B

raised by robotics startups by June 2026. All of 2025: $15B.

54–62%

of costs at leading AI companies go to compute.

30%+

of GenAI projects predicted to stop after proof of concept. Cost is a key reason.

THE LEADER VS. LAGGARD GAP · ELITE ENGINEERING TEAMS VS. LOW PERFORMERS

182× more deployments

127× faster from code to production

2,293× faster recovery from failure

Sources: Crunchbase (Jan 2026; Jun 2026) · Epoch AI (Feb 2026) · Gartner (Jul 2024) · DORA State of DevOps 2024

THE PROBLEM

But most teams are stuck between three bad options.

OPTION A

Hire a platform team

You get full control.

But:

- Senior platform and MLOps hires take months.
- You pay full salaries before anything ships.
- One key hire becomes a single point of failure.

OPTION B

Let your ML and robotics engineers do it

You add no cost on paper.

But:

- Your best people work on YAML, not models.
- Systems break in front of customers.
- GPU bills arrive as a surprise.

OPTION C

Hire a big consultancy or MSP

You get a thorough programme.

But:

- Slow to start and expensive.
- Seniors sell the work. Juniors deliver it.
- Few know GPUs, ROS 2 or robot fleets.

Every option costs you speed, money or control.

THE BETTER WAY

A senior platform team, on demand.

1 · 2 WEEKS

Diagnose

We map your cloud, cost and risks. You get a ranked plan.

2 · 4–6 WEEKS

Build

Fixed-scope sprints. We build the platform with your team.

3 · YOUR CHOICE

Hand over or operate

Your team takes over, or we stay on a monthly plan.

One architect with 20+ years of production experience works inside your team. Everything is code. Your team owns it at the end.

WHY THIS IS BETTER

Why this beats the alternatives

	DevLoops	Hire in-house	Do it yourself	Big consultancy
Time to start	Days	Months to hire	Now, but slow	Weeks to contract
Who does the work	20+ years, hands-on	Depends on the hire	ML team learns on the job	Juniors, after seniors sell
Cost model	Fixed price per phase	Ongoing salary and benefits	Hidden: lost product time	Open-ended day rates
AI and robotics depth	AV, robot and AI platforms	Hard to find locally	Models yes, infra no	Mostly generic cloud
After the work	Your team owns code and runbooks	Knowledge leaves with the person	Lives in one head	Often locked in

DIFFERENTIATORS

Are we a good fit for you?

Autonomy in production

Built cloud for an autonomous-vehicle platform and robot-swarm infrastructure.

Robotics depth

MSc Robotics. Builds ROS 2 fleet observability with OpenTelemetry.

Reliability record

100% uptime for 1.5+ years on a live AI learning platform.

Cost discipline

FinOps for cloud and GPU spend, from the first sprint.

Good fit if you:

- are an AI or robotics startup, seed to Series B
- take models or robots to production this year
- see cloud or GPU bills grow faster than revenue
- have no dedicated platform team yet

Not a fit if you:

- need a large 24/7 outsourced operations team
- want the lowest day rate for ticket work

PROOF

Do not take our word for it.

CLIENT

NeoHPC engineering team

[Result in one line: what you delivered and what changed for them.]

“[Testimonial: the result, in the client’s own words.]”

[Name, Title, Company]

“[Testimonial: the result, in the client’s own words.]”

[Name, Title, Company]

LIVE DEMO · 15 MIN

Trace one robot mission, end to end.

From ROS 2 on the robot, across the network, to the cloud. One trace in OpenTelemetry shows where time and failures go.

Or, on your own system: a live review of your cloud bill and cluster. You leave with three findings.

THE OFFER

How we work together

START HERE · 2 WEEKS

Diagnostic

Fixed price. You get:

- Architecture and security review
- Cloud and GPU cost review
- Reliability risks, ranked
- 90-day roadmap, with effort

THEN · 4–6 WEEKS

Build sprint

Fixed scope. For example:

- CI/CD and Infrastructure as Code
- Kubernetes and GPU platform
- MLOps or fleet observability
- Code, runbooks and team training

OPTIONAL · MONTHLY

Operate

Monthly plan. You get:

- Fractional architect, set days
- Support when incidents occur
- Monthly cost and reliability review
- Stop with 30 days’ notice

Fixed price for each phase, agreed before we start. The diagnostic fee counts toward the first sprint.

ONE QUESTION

If the diagnostic shows a clear path to production — and savings that cover its cost — will you start the first sprint with us?

Next step: a 30-minute call this week.

contact@devloopsconsulting.com · devloopsconsulting.com

Scan the code to pick a time.